

# Define model and import dependencies

```
In [ ]: import torch
        from torch import optim, nn, utils, Tensor
        import torch.nn.functional as F
        from torch.utils.data import DataLoader
        from torch.utils.data import Dataset
        import pandas as pd
        import numpy as np
```

```
In [ ]: # Defining our model
        class LSTMModel(nn.Module):
            def __init__(self, input_size, hidden_size, num_layers, num_classes):
                super(LSTMModel, self).__init__()
                self.hidden_size = hidden_size
                self.num_layers = num_layers
                self.lstm = nn.LSTM(input_size, hidden_size, num_layers, batch_first=True)
                self.fc = nn.Linear(hidden_size, num_classes)

            def forward(self, x):
                h0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(device=x.device)
                c0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(device=x.device)

                out, _ = self.lstm(x, (h0, c0))

                out = self.fc(out[:, -1, :])

                return out
```

```
In [ ]: from google.colab import drive
        drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [ ]: # Check if GPU is available. But should work fine on CPU too.
        if torch.cuda.is_available():
            device = torch.device('cuda')
        else:
            device = torch.device('cpu')

        print('device:', device)
```

device: cuda

## If you're training the model:

```
In [ ]: # train and validation data
        !unzip -q /content/drive/MyDrive/ybigta/2023-2/win_prediction_model/initial.zip -d /content

        # Open data_label, contains keypoints numpy array names and labels
        data_label = pd.read_csv("/content/initial/data_label.csv", index_col = 0)
```

```
In [ ]: # custom defined dataset for our data format
        class KeypointsDataset(Dataset):
```

```

def __init__(self, array_paths, labels, num_input):
    """
    Generates a custom numpy dataset

    Args:
        array_paths (list): list of numpy array inputs
        labels (list): list of labels
        num_input (int) : number of inputs to enter
    """
    self.array_paths = array_paths
    self.labels = labels
    self.num_input = num_input

def __len__(self):
    """
    Returns length of entire dataset
    """
    return len(self.array_paths)

def __getitem__(self, idx, ):
    """
    Gets the sample that corresponds to the sample id (idx)

    Args:
        idx (int): sample index

    Returns:
        keypoints (torch.Tensor): keypoints input tensor
        label (torch.Tensor): winning label tensor
    """
    keypoints_path = "/content/initial/data/"
    # load and turn numpy arrays into torch tensors
    keypoints = np.load(f"{keypoints_path}{self.array_paths[idx]}.npy")
    keypoints = keypoints[:, -self.num_input:]
    keypoints = torch.tensor(keypoints, dtype = torch.float32)
    label = torch.tensor(self.labels[idx])

    return keypoints, label

```

```

In [ ]: from sklearn.model_selection import train_test_split
        train_array, val_array = train_test_split(data_label, test_size = 0.1, random_state = 1

```

```

In [ ]: # Initialise the dataset and dataloader
        num_in = 100
        train_data = KeypointsDataset(
            array_paths = train_array['modified_json_filename'].reset_index(drop = True),
            labels = train_array['label'].reset_index(drop=True),
            num_input = num_in
        )
        val_data = KeypointsDataset(
            array_paths = val_array['modified_json_filename'].reset_index(drop= True),
            labels = val_array['label'].reset_index(drop= True),
            num_input = num_in
        )
        train_loader = DataLoader(train_data)
        val_loader = DataLoader(val_data)

```

```

In [ ]: def train_loop(dataloader, model, loss_fn, optimizer):
        size = len(dataloader.dataset)

```

```

# Set the model to training mode - important for batch normalization and dropout layers
# Unnecessary in this situation but added for best practices
model.train()
for i, (inputs, labels) in enumerate(dataloader):
    inputs = inputs.to(device)
    labels = labels.to(device)

    optimizer.zero_grad()

    outputs = model(inputs)

    loss = loss_fn(outputs, labels)
    loss.backward()

    optimizer.step()

    if i % 51 == 0:
        loss, current = loss.item(), (i + 1) * len(inputs)
        print(f"loss: {loss:>7f} [{current:>5d}/{size:>5d}]")

def test_loop(dataloader, model, loss_fn):
    # Set the model to evaluation mode - important for batch normalization and dropout layers
    # Unnecessary in this situation but added for best practices
    model.eval()
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    test_loss, correct = 0, 0

    # Evaluating the model with torch.no_grad() ensures that no gradients are computed
    # also serves to reduce unnecessary gradient computations and memory usage for tensors
    with torch.no_grad():
        for i, (inputs, labels) in enumerate(dataloader):
            inputs = inputs.to(device)
            labels = labels.to(device)

            outputs = model(inputs)

            _, predicted = torch.max(outputs.data, 1)
            test_loss += loss_fn(outputs, labels).item()
            correct += (outputs.argmax(1) == labels).type(torch.float).sum().item()
    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: {test_loss:>8f}")

```

```

In [ ]: # Setting up the model and initiating the train & validation loops
model = LSTMModel(input_size=num_in, hidden_size=10, num_layers=1, num_classes=2).to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.0001, weight_decay = 0.1)

num_epochs = 60
for t in range(num_epochs):
    print(f"Epoch {t+1}\n-----")
    train_loop(train_loader, model, criterion, optimizer)
    test_loop(val_loader, model, criterion)
print("Done!")

```

Epoch 1

---

loss: 0.753100 [ 1/ 54]  
loss: 0.844226 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.737354

Epoch 2

---

loss: 0.754023 [ 1/ 54]  
loss: 0.845346 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.735843

Epoch 3

---

loss: 0.755222 [ 1/ 54]  
loss: 0.846515 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.734260

Epoch 4

---

loss: 0.756355 [ 1/ 54]  
loss: 0.847704 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.732601

Epoch 5

---

loss: 0.757482 [ 1/ 54]  
loss: 0.848895 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.730864

Epoch 6

---

loss: 0.758621 [ 1/ 54]  
loss: 0.850077 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.729048

Epoch 7

---

loss: 0.759782 [ 1/ 54]  
loss: 0.851233 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.727152

Epoch 8

---

loss: 0.760967 [ 1/ 54]  
loss: 0.852350 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.725174

Epoch 9

---

loss: 0.762177 [ 1/ 54]  
loss: 0.853411 [ 52/ 54]

Test Error:  
Accuracy: 33.3%, Avg loss: 0.723113

Epoch 10

---

loss: 0.763412 [ 1/ 54]  
loss: 0.854398 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.720966

Epoch 11

---

loss: 0.764672 [ 1/ 54]  
loss: 0.855291 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.718730

Epoch 12

---

loss: 0.765954 [ 1/ 54]  
loss: 0.856072 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.716403

Epoch 13

---

loss: 0.767257 [ 1/ 54]  
loss: 0.856721 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.713981

Epoch 14

---

loss: 0.768579 [ 1/ 54]  
loss: 0.857220 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.711463

Epoch 15

---

loss: 0.769916 [ 1/ 54]  
loss: 0.857551 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.708845

Epoch 16

---

loss: 0.771267 [ 1/ 54]  
loss: 0.857699 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.706126

Epoch 17

---

loss: 0.772627 [ 1/ 54]  
loss: 0.857655 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.703305

Epoch 18

---

loss: 0.773994 [ 1/ 54]  
loss: 0.857410 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.700380

Epoch 19

---

loss: 0.775363 [ 1/ 54]  
loss: 0.856965 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.697349

Epoch 20

---

loss: 0.776730 [ 1/ 54]  
loss: 0.856322 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.694207

Epoch 21

---

loss: 0.778091 [ 1/ 54]  
loss: 0.855491 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.690947

Epoch 22

---

loss: 0.779440 [ 1/ 54]  
loss: 0.854485 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.687556

Epoch 23

---

loss: 0.780770 [ 1/ 54]  
loss: 0.853324 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.684013

Epoch 24

---

loss: 0.782076 [ 1/ 54]  
loss: 0.852027 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.680290

Epoch 25

---

loss: 0.783349 [ 1/ 54]  
loss: 0.850616 [ 52/ 54]  
Test Error:  
Accuracy: 33.3%, Avg loss: 0.676348

Epoch 26

---

loss: 0.784582 [ 1/ 54]  
loss: 0.849112 [ 52/ 54]  
Test Error:

Accuracy: 33.3%, Avg loss: 0.672157

Epoch 27

-----  
loss: 0.785767 [ 1/ 54]

loss: 0.847522 [ 52/ 54]

Test Error:

Accuracy: 33.3%, Avg loss: 0.667716

Epoch 28

-----  
loss: 0.786895 [ 1/ 54]

loss: 0.845837 [ 52/ 54]

Test Error:

Accuracy: 33.3%, Avg loss: 0.663079

Epoch 29

-----  
loss: 0.787960 [ 1/ 54]

loss: 0.844027 [ 52/ 54]

Test Error:

Accuracy: 33.3%, Avg loss: 0.658367

Epoch 30

-----  
loss: 0.788960 [ 1/ 54]

loss: 0.842068 [ 52/ 54]

Test Error:

Accuracy: 33.3%, Avg loss: 0.653736

Epoch 31

-----  
loss: 0.789901 [ 1/ 54]

loss: 0.839984 [ 52/ 54]

Test Error:

Accuracy: 33.3%, Avg loss: 0.649314

Epoch 32

-----  
loss: 0.790797 [ 1/ 54]

loss: 0.837869 [ 52/ 54]

Test Error:

Accuracy: 33.3%, Avg loss: 0.645132

Epoch 33

-----  
loss: 0.791672 [ 1/ 54]

loss: 0.835827 [ 52/ 54]

Test Error:

Accuracy: 50.0%, Avg loss: 0.641122

Epoch 34

-----  
loss: 0.792546 [ 1/ 54]

loss: 0.833888 [ 52/ 54]

Test Error:

Accuracy: 50.0%, Avg loss: 0.637193

Epoch 35

loss: 0.793424 [ 1/ 54]  
loss: 0.832006 [ 52/ 54]  
Test Error:  
Accuracy: 50.0%, Avg loss: 0.633305

Epoch 36

-----  
loss: 0.794302 [ 1/ 54]  
loss: 0.830107 [ 52/ 54]  
Test Error:  
Accuracy: 50.0%, Avg loss: 0.629461

Epoch 37

-----  
loss: 0.795168 [ 1/ 54]  
loss: 0.828129 [ 52/ 54]  
Test Error:  
Accuracy: 50.0%, Avg loss: 0.625678

Epoch 38

-----  
loss: 0.796018 [ 1/ 54]  
loss: 0.826027 [ 52/ 54]  
Test Error:  
Accuracy: 50.0%, Avg loss: 0.621969

Epoch 39

-----  
loss: 0.796846 [ 1/ 54]  
loss: 0.823774 [ 52/ 54]  
Test Error:  
Accuracy: 50.0%, Avg loss: 0.618336

Epoch 40

-----  
loss: 0.797650 [ 1/ 54]  
loss: 0.821353 [ 52/ 54]  
Test Error:  
Accuracy: 50.0%, Avg loss: 0.614779

Epoch 41

-----  
loss: 0.798430 [ 1/ 54]  
loss: 0.818757 [ 52/ 54]  
Test Error:  
Accuracy: 50.0%, Avg loss: 0.611292

Epoch 42

-----  
loss: 0.799186 [ 1/ 54]  
loss: 0.815979 [ 52/ 54]  
Test Error:  
Accuracy: 50.0%, Avg loss: 0.607869

Epoch 43

-----  
loss: 0.799919 [ 1/ 54]  
loss: 0.813020 [ 52/ 54]  
Test Error:  
Accuracy: 66.7%, Avg loss: 0.604503



Epoch 44

---

loss: 0.800629 [ 1/ 54]  
loss: 0.809878 [ 52/ 54]  
Test Error:  
Accuracy: 66.7%, Avg loss: 0.601188

Epoch 45

---

loss: 0.801317 [ 1/ 54]  
loss: 0.806558 [ 52/ 54]  
Test Error:  
Accuracy: 66.7%, Avg loss: 0.597916

Epoch 46

---

loss: 0.801984 [ 1/ 54]  
loss: 0.803063 [ 52/ 54]  
Test Error:  
Accuracy: 66.7%, Avg loss: 0.594682

Epoch 47

---

loss: 0.802630 [ 1/ 54]  
loss: 0.799401 [ 52/ 54]  
Test Error:  
Accuracy: 66.7%, Avg loss: 0.591482

Epoch 48

---

loss: 0.803257 [ 1/ 54]  
loss: 0.795585 [ 52/ 54]  
Test Error:  
Accuracy: 66.7%, Avg loss: 0.588311

Epoch 49

---

loss: 0.803864 [ 1/ 54]  
loss: 0.791629 [ 52/ 54]  
Test Error:  
Accuracy: 66.7%, Avg loss: 0.585167

Epoch 50

---

loss: 0.804453 [ 1/ 54]  
loss: 0.787553 [ 52/ 54]  
Test Error:  
Accuracy: 66.7%, Avg loss: 0.582049

Epoch 51

---

loss: 0.805022 [ 1/ 54]  
loss: 0.783378 [ 52/ 54]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.578955

Epoch 52

---

loss: 0.805574 [ 1/ 54]

loss: 0.779128 [ 52/ 54]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.575886

Epoch 53

---

loss: 0.806108 [ 1/ 54]  
loss: 0.774824 [ 52/ 54]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.572842

Epoch 54

---

loss: 0.806623 [ 1/ 54]  
loss: 0.770486 [ 52/ 54]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.569824

Epoch 55

---

loss: 0.807120 [ 1/ 54]  
loss: 0.766131 [ 52/ 54]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.566831

Epoch 56

---

loss: 0.807598 [ 1/ 54]  
loss: 0.761770 [ 52/ 54]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.563865

Epoch 57

---

loss: 0.808056 [ 1/ 54]  
loss: 0.757412 [ 52/ 54]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.560926

Epoch 58

---

loss: 0.808496 [ 1/ 54]  
loss: 0.753059 [ 52/ 54]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.558014

Epoch 59

---

loss: 0.808915 [ 1/ 54]  
loss: 0.748713 [ 52/ 54]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.555128

Epoch 60

---

loss: 0.809313 [ 1/ 54]  
loss: 0.744372 [ 52/ 54]  
Test Error:  
Accuracy: 83.3%, Avg loss: 0.552268

Done!

```
In [ ]: # If you're continuing to train from a pre-trained version:
pretrained_path = "/content/drive/MyDrive/ybigta/2023-2/win_prediction_model/winpred_model.pth"
model = LSTMModel(input_size=num_in, hidden_size=10, num_layers=1, num_classes=2).to(device)
model.load_state_dict(torch.load(pretrained_path, map_location=torch.device("cpu")), strict=True)

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.0001, weight_decay = 0.1)

num_epochs = 5
for t in range(num_epochs):
    print(f"Epoch {t+1}Wn-----")
    train_loop(train_loader, model, criterion, optimizer)
    test_loop(val_loader, model, criterion)
print("Done!")
```

Epoch 1

```
-----
loss: 0.783462 [ 1/ 54]
loss: 0.599811 [ 52/ 54]
Test Error:
Accuracy: 83.3%, Avg loss: 0.532518
```

Epoch 2

```
-----
loss: 0.779629 [ 1/ 54]
loss: 0.564055 [ 52/ 54]
Test Error:
Accuracy: 83.3%, Avg loss: 0.524813
```

Epoch 3

```
-----
loss: 0.777854 [ 1/ 54]
loss: 0.543219 [ 52/ 54]
Test Error:
Accuracy: 83.3%, Avg loss: 0.520422
```

Epoch 4

```
-----
loss: 0.776490 [ 1/ 54]
loss: 0.530013 [ 52/ 54]
Test Error:
Accuracy: 83.3%, Avg loss: 0.517872
```

Epoch 5

```
-----
loss: 0.775345 [ 1/ 54]
loss: 0.520610 [ 52/ 54]
Test Error:
Accuracy: 83.3%, Avg loss: 0.516346
```

Done!

## If you're doing a test:

### Test Dataset Version

```
In [ ]: !unzip -q /content/drive/MyDrive/ybigta/2023-2/win_prediction_model/test.zip -d /content
```

```
In [ ]: PATH_TO_NUMPY_ARRAY_FOLDER = "/content/test/data/"
PATH_TO_NUMPY_ARRAY_FILENAMES_AND_LABELS = "/content/test/test_label.csv"
PATH_TO_PRETRAINED_MODEL_STATE_DICT = "/content/drive/MyDrive/ybigta/2023-2/win_predict
```

```
In [ ]: class KeypointsDataset(Dataset):
    def __init__(self, array_paths, labels, num_input):
        """
        Generates a custom numpy dataset

        Args:
            array_paths (list): list of numpy array inputs
            labels (list): list of labels
            num_input (int) : number of inputs to enter
        """
        self.array_paths = array_paths
        self.labels = labels
        self.num_input = num_input

    def __len__(self):
        """
        Returns length of entire dataset
        """
        return len(self.array_paths)

    def __getitem__(self, idx, ):
        """
        Gets the sample that corresponds to the sample id (idx)

        Args:
            idx (int): sample index

        Returns:
            keypoints (torch.Tensor): keypoints input tensor
            label (torch.Tensor): winning label tensor
        """
        keypoints_path = PATH_TO_NUMPY_ARRAY_FOLDER
        # Turn numpy arrays into torch tensors
        keypoints = np.load(f"{keypoints_path}{self.array_paths[idx]}.npy")
        keypoints = keypoints[:, -self.num_input:]
        keypoints = torch.tensor(keypoints, dtype = torch.float32)
        label = torch.tensor(self.labels[idx])

        return keypoints, label

def test_loop(dataloader, model, loss_fn):
    # Set the model to evaluation mode - important for batch normalization and dropout
    # Unnecessary in this situation but added for best practices
    model.eval()
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    test_loss, correct = 0, 0

    # Evaluating the model with torch.no_grad() ensures that no gradients are computed
    # also serves to reduce unnecessary gradient computations and memory usage for tensors
    with torch.no_grad():
        for i, (inputs, labels) in enumerate(dataloader):
            inputs = inputs.to(device)
```

```

        labels = labels.to(device)

        outputs = model(inputs)

        _, predicted = torch.max(outputs.data, 1)
        test_loss += loss_fn(outputs, labels).item()
        correct += (outputs.argmax(1) == labels).type(torch.float).sum().item()
    test_loss /= num_batches
    correct /= size
    print(f"Test Error: Wn Accuracy: {(100*correct):>0.1f}%, Avg loss: {test_loss:>8f}")

test_array = pd.read_csv(PATH_TO_NUMPY_ARRAY_FILENAMES_AND_LABELS)
num_in = 100 # DO NOT ALTER
test_data = KeypointsDataset(
    array_paths = test_array['modified_json_filename'].reset_index(drop=True),
    labels = test_array['label'].reset_index(drop=True),
    num_input = num_in
)
test_loader = DataLoader(test_data)
pretrained_path = PATH_TO_PRETRAINED_MODEL_STATE_DICT
model = LSTMModel(input_size=num_in, hidden_size=10, num_layers=1, num_classes=2).to(device)
model.load_state_dict(torch.load(pretrained_path, map_location=torch.device("cuda")),
criterion = nn.CrossEntropyLoss())

```

```

In [ ]: # Run this after setting all the paths, and the above cell (loading the model)
        test_loop(test_loader, model, criterion)

```

```

Test Error:
Accuracy: 50.0%, Avg loss: 0.715470

```

```

In [ ]: print(model)

LSTMModel(
  (lstm): LSTM(100, 10, batch_first=True)
  (fc): Linear(in_features=10, out_features=2, bias=True)
)

```

## Individual File Version

```

In [ ]: !unzip -q /content/drive/MyDrive/ybigta/2023-2/win_prediction_model/test_keypoints.zip
        replace /content/test_keypoints/31_person.npy? [y]es, [n]o, [A]ll, [N]one, [r]ename: A

```

```

In [ ]: PATH_TO_NUMPY_ARRAY = "/content/test/data/lee_kim_1_person.npy"
        PATH_TO_PRETRAINED_MODEL_STATE_DICT = "/content/drive/MyDrive/ybigta/2023-2/win_predict

```

```

In [ ]: numpy_data = np.load(PATH_TO_NUMPY_ARRAY)

        pretrained_path = PATH_TO_PRETRAINED_MODEL_STATE_DICT
        model = LSTMModel(input_size=num_in, hidden_size=10, num_layers=1, num_classes=2).to(device)
        # If there's no GPU, change from device("cuda") to device("cpu")
        model.load_state_dict(torch.load(pretrained_path, map_location=torch.device("cuda")),

        def test(numpy_data_raw, model):
            data = torch.tensor(numpy_data_raw[:, -100:], dtype=torch.float32)
            data = data[None, :]
            data = data.to(device)
            model.eval()
            outputs = model(data)

```

```
_, predicted = torch.max(outputs.data, 1)
predicted_int = predicted.type(torch.int)[0]
print(f"The result is:{predicted_int:d}")
return predicted_int
test(numpy_data, model)
```

The result is:0

out[ ]: tensor(0, device='cuda:0', dtype=torch.int32)